



Implementing Version Control With Git and GitHub as a Learning Objective in Statistics and Data Science Courses

Matthew D. Beckman, Mine Çetinkaya-Rundel, Nicholas J. Horton, Colin W. Rundel, Adam J. Sullivan & Maria Tackett

To cite this article: Matthew D. Beckman, Mine Çetinkaya-Rundel, Nicholas J. Horton, Colin W. Rundel, Adam J. Sullivan & Maria Tackett (2021) Implementing Version Control With Git and GitHub as a Learning Objective in Statistics and Data Science Courses, Journal of Statistics and Data Science Education, 29:sup1, S132-S144, DOI: [10.1080/10691898.2020.1848485](https://doi.org/10.1080/10691898.2020.1848485)

To link to this article: <https://doi.org/10.1080/10691898.2020.1848485>



© 2021 The Author(s). Published with license by Taylor & Francis Group, LLC.



Published online: 22 Mar 2021.



Submit your article to this journal [↗](#)



Article views: 8017





View related articles [↗](#)



Citing articles: 7 View citing articles [↗](#)

Implementing Version Control With Git and GitHub as a Learning Objective in Statistics and Data Science Courses

Matthew D. Beckman^a, Mine Çetinkaya-Rundel^{b,c,d} , Nicholas J. Horton^e , Colin W. Rundel^{b,d}, Adam J. Sullivan^f, and Maria Tackett^d

^aDepartment of Statistics, Penn State University, State College, PA; ^bSchool of Mathematics, University of Edinburgh, Edinburgh, UK; ^cRStudio, Boston, MA; ^dDepartment of Statistical Science, Duke University, Durham, NC; ^eDepartment of Mathematics and Statistics, Amherst College, Amherst, MA; ^fDepartment of Biostatistics, Brown University, Providence, RI

ABSTRACT

A version control system records changes to a file or set of files over time so that changes can be tracked and specific versions of a file can be recalled later. As such, it is an essential element of a reproducible workflow that deserves due consideration among the learning objectives of statistics courses. This article describes experiences and implementation decisions of four contributing faculty who are teaching different courses at a variety of institutions. Each of these faculty has set version control as a learning objective and successfully integrated one such system (Git) into one or more statistics courses. The various approaches described in the article span different implementation strategies to suit student background, course type, software choices, and assessment practices. By presenting a wide range of approaches to teaching Git, the article aims to serve as a resource for statistics and data science instructors teaching courses at any level within an undergraduate or graduate curriculum.

KEYWORDS

Collaborative learning; Data acumen; Education; Reproducible analysis; Statistical computing; Workflow

1. Introduction

Nolan and Temple Lang (2010) promoted “version control” as a key topic for statistical analysis, particularly when coordinating work across a team. A version control system records changes to a file or set of files over time so that changes can be tracked and specific versions of a file can be recalled later.

The 2014 *American Statistical Association Curriculum Guidelines for Undergraduate Programs* includes proficiency with modern statistical software as well as well-documented and reproducible data wrangling skills as a necessary component of the undergraduate statistics curriculum (American Statistical Association 2014). The National Academies consensus report on *Data Science for Undergraduates* (National Academies of Science, Engineering, and Medicine 2018) identifies workflow and reproducibility as important components of “data acumen.” Version control is an important foundation for reproducible workflows, be they collaborative (maintaining versions of files that are being modified by teams) or non-collaborative (tracking analysis histories and providing analysis provenance). It forms a necessary part of a reproducible workflow, and therefore deserves due consideration among the learning objectives of statistics and data science courses. Fiksel et al. (2019) motivated the use of GitHub for version control and describe how they integrated this complex and powerful system into two courses.

This article follows a similar format to Garfield et al. (2011) and Hardin et al. (2015) by describing the experiences and implementation decisions of several contributing faculty—teaching different courses at different institutions—who have

successfully integrated Git into one or more statistics courses to teach version control as a learning objective. We begin by discussing our motivations for identifying version control as a learning objective and then provide summaries of courses taught by the four contributing faculty highlighting different implementation strategies chosen based on student audience, course type, software choices, and assessment practices. We highlight a range of implementations across a variety of courses and student populations to provide a resource for statistics instructors to interpolate an implementation suitable for use in their own courses at the undergraduate or graduate level. We refer the reader to Table 1 for definitions of terms we will use regularly throughout the article. Readers who are unfamiliar with version control would benefit from reading (Bryan 2018a).

1.1. Motivation for Version Control

There are two main motivations for including version control as a learning objective in statistics courses. The first motivation is reproducibility. For a scientific study to be replicated, the statistical analysis in the study must be entirely reproducible. Teaching reproducible analysis in the statistics curriculum helps make students aware of the issue of scientific reproducibility and also equips them with the knowledge and skills to conduct their future data analyses reproducibly, whether as part of an academic research project or in industry. Baumer et al. (2014) advocated teaching literate programming early in the statistics curriculum via the use of R Markdown, a system that enables stu-

Table 1. Definitions of common terms.

Term	Definition
Git	An open source version control software system (git-scm.com)
Git repository (or repo)	Analogous to a project directory location or a folder in Google Drive, Dropbox, etc. It tracks changes to files.
GitHub	A remote commercial hosting service for Git repositories (GitHub 2020a)
GitHub issues	A mechanism to track tasks or ideas
commit	A set of saved changes to a local repo
pull	Update a local repo
push	Upload local files to a remote repo
forking	Create a copy of a repository under your account
pull request	Propose changes to a remote repo
merge conflict	Contradictory changes that cannot be integrated until they are reconciled by a user
branching	Keeping multiple snapshots of a repo
gh-pages (GitHub Pages)	Special branch which allows creation of a webpage from within GitHub
GitHub Actions	Mechanism for continuous integration
GitHub Classroom	A system to facilitate distributing assignments to students. Instructors create a template Git repository that includes starter code, datasets, and document templates that students may need. A single URL is provided to the class, and each student is provided their own copy of the template repository when they click the URL and accept the assignment. The instructor can reuse the template repositories in future offerings (GitHub Education 2020).
ghclass	An R package that provides an alternative system to GitHub Classroom to facilitate distributing assignments to students (Rundel, Çetinkaya-Rundel, and Anders 2020).
RStudio	An Integrated Development Environment (IDE), that is, a front-end, for R that offers integration with Git. (rstudio.com)
RStudio Server Pro	A server-based version of RStudio that can be installed for free for academic use by instructors or institutions. (rstudio.com/products/rstudio-server-pro)
RStudio Cloud	A cloud-based version of RStudio software on servers provisioned by RStudio. (rstudio.cloud)

dents to produce computational documents that include their code, output, and written analysis using the `rmarkdown` package (Xie, Allaire, and Golemund 2018). Literate programming with R Markdown goes a long way toward computational reproducibility, but a data analysis of considerable scope likely cannot be managed in a single R Markdown document. As Bryan (2018a) puts it, data analysis is an iterative process that relies on and produces many files—input data, source code, figures, tables, reports, etc. Managing such projects is not unique to statistics, but it is something that our curricula have been slow to address. Version control provides a mechanism for managing all these files and sharing them with others as a project progresses, and modern tooling and workflows make it easier to implement in teaching than ever before.

The second motivation is industry and academic preparedness. The ability to use version control systems is a highly desired skill in any industry where writing code is part of the job, and the need to teach it has been recognized in the literature (Haaranen and Lehtinen 2015). Git is a widely used tool in industry for version control and code sharing. In a 2017 survey of data scientists conducted by Kaggle, over 58% of 6000 survey respondents remarked that Git was the main system used for version control and code sharing in their workplace (Kaggle 2017). Additionally, knowing how to use GitHub is considered an essential skill in the tech field, just as important as software development and technical writing (Zagalsky et al. 2015). In an era where many of our statistics and data science students are heading into jobs where they will be writing code and working alongside software engineers and developers, it is essential that we equip them with these skills.

Exposure to version control early (and often) in a statistics curriculum ensures that, by the end, students not only enhance their statistical analysis skills, but also develop workflows for conducting analyses individually and collaboratively. With the widespread use of GitHub in academia and industry, courses that teach version control prepare students for internships,

research programs, and their future careers. More immediately, they can use these computing tools as they work on analyses and projects in subsequent courses.

Additionally, implementing version control in a course can encourage students to think about statistical analysis as an iterative process. While working on a given assignment, students “submit” their work multiple times by knitting their R Markdown file, writing a commit message to document the changes, and pushing the updated work to their assignment repository (repo). Some have adopted the mantra “knit-commit-push” for this workflow, and others “commit-pull-push.” Both are effective ways to help drive home the way that GitHub structures and organizes changes to files.

Use of version control helps reinforce the notion that statistical analysis typically requires multiple revisions, as students can review their commits to see all the updates they’ve made to their work. A desirable side effect is that, because students are periodically “submitting” their assignment as they work on it, there is less pressure of the final deadline where everything must be submitted in its final form. By the time a deadline approaches, students have ideally submitted a majority of their work, which may reduce issues around late submissions.

2. Method

To organize this article, the authors first agreed upon a set of organizing prompts to provide direction as they describe their experiences:

- Describe the course/students.
- Why use Git and GitHub?
- What tools do you use for implementing Git in your class?
- How do you introduce Git? Describe your students’ first encounter with Git in the course.
- What role does Git have in the regular day-to-day workflow for your students in the course?

- How do you assess Git proficiency as a learning objective?
- How do you address the United States Family Educational Rights and Privacy Act (FERPA) and related privacy issues?
- Do you have other advice for instructors considering incorporating Git or some other type of version control into a statistics course?

The contributors were then free to address as many of these prompts as they deemed appropriate. Each narrative description was then written independently in an attempt to reduce cross-pollination and promote similarities and differences to emerge naturally. The panel responses have been organized to follow a similar structure within each section (course description, tools and implementation, first exposure in class, regular workflow, assessment, and other remarks). The order in which these responses are presented aligns roughly with their place in a statistics curriculum at each respective institution: a first year undergraduate course, a second undergraduate course in statistics, a course in a Master's in Statistical Science program, and finally a Master's level course in a Biostatistics program.

3. Common Features of the Courses

While taught at different levels and serving different audiences, there are also a few aspects shared by all of the courses described. First, all of these courses teach and use the R computing language along with RStudio as the integrated development environment for R (RStudio Team 2015). Second, each course either requires or offers the option to access RStudio in the browser, either using an RStudio Server Pro instance hosted by their university or using RStudio Cloud, a cloud-based service managed by RStudio (RStudio Team 2020). Both options allow students and the instructor to use the same versions of R, RStudio, and any packages required for the course, which cuts down on early difficulties related to managing local installations. Additionally, this means students only need a web browser and an internet connection to access the computing environment and can start programming as soon as the first day of class (Çetinkaya-Rundel and Rundel 2018). Nearly any computer, Chromebook, or tablet is sufficient, and students can easily switch hardware as needed if using one device (e.g., tablet or lab computer) in class and

another (e.g., personal computer) outside of class. In the rest of the article, we will refer to these products generically as RStudio or the RStudio IDE.

Server-based access to RStudio also streamlines Git installation and integration with RStudio. Each course uses the Git pane in RStudio as opposed to Git's command line interface or GitHub Desktop. The RStudio interface is attractive since it is familiar to many students and it facilitates use of the basic functionality of Git using a point-and-click interface to practice version control fundamentals and implement key steps of the workflow (e.g., *diff*, *commit*, *pull*, *push*; see Figure 1). This implementation serves to mitigate cognitive load while students gain proficiency with unfamiliar tools and workflows, yet easily extends when required, since the terminal is accessible within RStudio if shell commands are necessary.

All of the instructors manage Git related course logistics by establishing a GitHub Pro organization through GitHub Education. GitHub provides unlimited free private repositories as well as compute-time credits that can be used for running automated actions on the student repositories. Use of private repositories ensures compliance with the Federal Educational Rights and Privacy Act (FERPA) by protecting student work and information from public access (Federal Educational Rights and Privacy Act (FERPA) 2020). These private repositories are only accessible to the student, the instructor, and any other course teaching staff, for example, teaching assistants. Additionally, GitHub organizations hide the identity of all members to nonmembers by default, meaning that students' enrollment in the course will not be disclosed by their joining the course organization. Access within the GitHub organization is further controlled on a per user basis via permissions: the instructor and teaching assistants are "Owners" and the students are "Members." As "Owners," the instructor and teaching assistants are able to manage organization membership as well as create and manage any repository created within the organization, and thus see all students' work. As "Members," students are only able to view and access the individual or team repositories assigned to them; they cannot view or access the repositories for any other students. They, as well as nonmembers, can also view any repositories within the organization that have been made public, such as those containing supplemental notes or any other materials for the course.

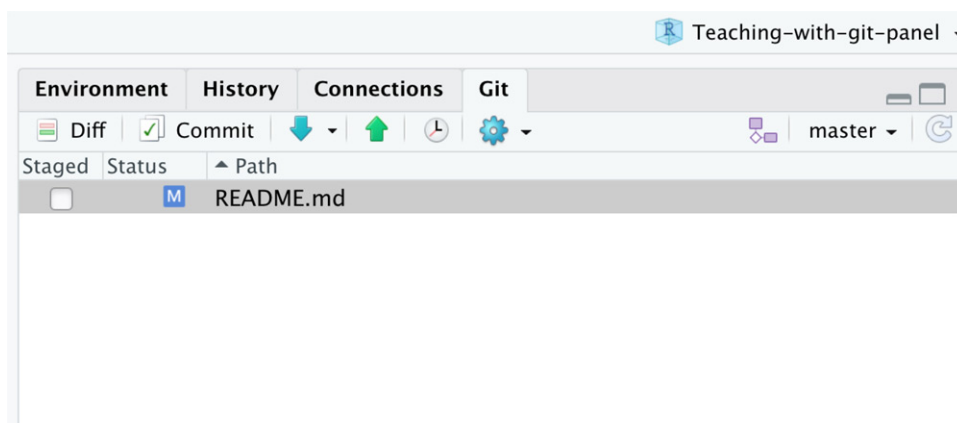


Figure 1. Example of Git pane within an RStudio IDE window.

4. First-Year Data Computing Course (Beckman)

4.1. Course Description

STAT 184: Introduction to R Programming at Penn State University is a two-credit (2×50 min instruction each week for 15 weeks) R programming course originally modeled after a similar course and accompanying textbook (Kaplan 2015; Kaplan and Beckman 2019) first developed by Daniel Kaplan at Macalester College. This course is designed for first-year undergraduate students from any academic program and has no prerequisites. The course currently enrolls 30–40 students in each of 9 sections per academic year, although enrollment demand has been increasing rapidly since it was first developed in 2015. At least one section each year is made available to first-semester students interested in the statistics major (it has even been coordinated with their orientation seminar class in the past) and at least one other section is offered by popular demand to mixed audiences of any major and class standing.

Major topics in STAT 184 include data wrangling and visualization with *tidyverse* tools, literate programming with R Markdown, and version control with Git. These themes persist for the entire semester and are complemented by a survey of topics such as statistical foundations, web scraping, regular expressions, simulation principles, and basic machine learning ideas.

4.2. Tools and Implementation

The workflow for students in STAT 184 includes the RStudio IDE, its Git pane for version control, and interacting with GitHub. The only step in the workflow that involves a tool outside of these is typically the Git configuration step that students complete once at the beginning of the semester using the Terminal pane of RStudio.

The instructors use one additional tool, GitHub Classroom (GitHub Education 2020), to deploy assignments to students as described in Fiksel et al. (2019). GitHub Classroom facilitates the batch creation of private student repositories on GitHub with starter documents for each assignment such as instructions, a grading rubric, data sources, and an R Markdown template.

Each week students are assigned one or more assignments with GitHub Classroom. Students then clone these repositories, work on them and commit and push their changes back to GitHub as they go, and finally submit their assignments in the university's learning management system (LMS). The end-of-semester project has a slightly different workflow; work may be submitted as a GitHub repository or a website automatically created using GitHub Pages, a setting configured within the project repository.

4.3. First Exposure in Class

For the first student encounter with Git and GitHub in STAT 184, the instructor creates a public GitHub repository associated with a GitHub Pages website for the class (see, e.g., mdbeckman.github.io/GitHub-Practice-StatChat-SP20). After a 15-min class discussion to motivate version control as a means to support collaboration and reproducibility as well as orient students to

a schematic representation of a workflow that includes Git and GitHub, students are provided a link to the aforementioned GitHub Pages website which includes instructions for a hands-on activity to be completed during class. The activity walks through a first encounter with Git and GitHub that invites students to (1) create a GitHub profile of their own, (2) create their first Git repository, (3) turn their personal repository into a GitHub Pages website, (4) edit a table in the instructor's class repository to add their name, GitHub user ID, and a functioning URL for GitHub Pages website they have just created, (5) leave an informative commit message and initiate a *pull request* to the instructor's repository. The resulting table (from step 3) provides the instructor with the name and GitHub username associated with each student, and every student creates a public GitHub Pages website as a starting point to begin developing a work portfolio for future use.

The exercise takes approximately 30 min and introduces key elements of version control:

- Creating a repository.
- Making a few commits and issuing a pull request.
- Contributing to an outside repository belonging to another GitHub user.
- Observing how to merge pull requests.
- Observing how merge conflicts are created and resolved.

This short activity has the additional benefit that the entire exercise can be completed within the GitHub web interface. The advantage of this approach is that students do not need to use R, the Terminal, or even Markdown, which allows them to begin building a schema for version control without distraction from tools still unfamiliar to them at this early stage.

4.4. Workflow

After the first exposure to version control as described above, each assignment throughout the semester is associated with a private GitHub repository that each student maintains. Nearly the entire workflow takes place within the RStudio IDE. The only regular (albeit trivial) version control task outside the RStudio IDE is the requirement that students click a link to accept the template repository deployed through GitHub Classroom, and then students are taken to a GitHub repository associated with their personal copy of the assignment which they clone as an RStudio project. From that point on, students can make a *commit*, *push*, *pull*, view a *diff* (difference between current and previous versions of a document), etc. using the Git pane in RStudio.

Some of these Git repositories are associated with an activity that is launched, completed, and submitted in the space of a week or less, while other Git repositories are used on a regular basis throughout nearly the entire academic term. Additionally, a few repositories are associated with collaborative assignments for which two or three students must contribute by making commits to shared repositories. Students are expected to make regular commits in each repository with the assessment of several assignments taking into account their commit history evident in the associated GitHub repos. However, the final product of most assignments is submitted to the course LMS for grading.

4.5. Assessment

If version control is to be taken seriously as a learning objective for a course, then it should be made clear to students early and often. Students should find mention in the course syllabus, students should expect to see it on exams, and students should feel that it is among the class norms for regular assignments. In the earliest iterations of STAT 184 development, version control had been treated as an incidental topic: encouraged, but not assessed. As a learning objective, version control is now integrated into a wide variety of assessments, including homework assignments, projects, and exams. For early assignments, success is set at a relatively low threshold: objective evidence that students have simply created a repository of their own or edited a repository provided to them. As the workflow becomes more familiar, assessment may include direct scrutiny of a commit history or similar activity documented within a specific repository.

For example, students are expected to maintain a single repository for all weekly problem sets assigned from the textbook during the semester. This repository is then graded as a distinct homework score at the end of the semester based on verification that all assignments are present and associated with some minimum number of commits per assignment. To be clear, the goal is simply to incentivize commits early and often in the workflow of each assignment. Students should not be preoccupied by counting commits; the actual number is inconsequential. Once version control has truly taken root in the students' workflow, many STAT 184 students more than double the minimum number of commits required.

Finally, students are to expect version control content on in-class exams. For example, this might include open-ended tasks about important concepts, selected-response questions (e.g., True/False or multiple choice) about procedural details such as whether a *pull* action modifies files (a) in the directory on their local computer [correct answer], (b) on the GitHub Remote Server, (c) both, (d) neither, or another task might prompt students to resolve an apparent merge conflict presented to them in a screenshot.

4.6. A Subsequent Data Science Course

STAT 184, along with an introductory statistics course (e.g., AP Statistics), serves as a prerequisite for an intermediate-level course in the curriculum (STAT 380: Data Science Through Statistical Reasoning and Computation) that is required for both statistics and data science majors. This course extends practices, including version control, introduced in STAT 184 such that the tools and workflow are largely unchanged with the exception that students are expected to use a local installation of RStudio. Since prior experience with version control in the prerequisite STAT 184 class is assumed, STAT 380 instructors can simply upload a roster of names and E-mail addresses to GitHub Classroom and students can identify themselves as they start their first assignment. Assessments still include various version control elements but are held to higher standards as expected of a more mature workflow.

5. A Second Course in Statistics (Tackett)

5.1. Course Description

STA 210: Regression Analysis is an intermediate-level course in the Department of Statistical Science at Duke University. About 90 students take the course each semester, representing a variety of majors across campus. The course is one of the core requirements for the statistical science major and minor, so a large proportion of the students intend to pursue the major or minor. The only prerequisite is an introductory statistics or probability course, therefore, the students coming into STA 210 have a range of previous experiences using R and Git. As an example, in Fall 2019, a majority of students had previously used R and RStudio in another course, while less than half had any previous exposure to Git and GitHub. Given the variability in previous experiences with these computing tools, some of the challenges of teaching Git in this course are similar to those experienced in a first semester statistics course.

5.2. Tools and Implementation

The primary computing tools used in STA 210 are the RStudio IDE and GitHub. On the instructor side, all administrative activities involving GitHub are done using the `ghclass` R package (Rundel, Çetinkaya-Rundel, and Anders 2020). These activities include adding the students to the GitHub organization at the beginning of the semester, creating teams on GitHub, making and replicating assignment repos, and cloning the repositories for grading. These processes are described in more detail in [Section 5.4](#) and further details are provided in the documentation for the `ghclass` package. The computing infrastructure using RStudio and GitHub as well as the course pedagogy is based on Çetinkaya-Rundel and Rundel (2018) and Data Science in a Box (Çetinkaya-Rundel 2020).

5.3. First Exposure in Class

Students are introduced to Git at the very beginning of the semester. On the first day of class, students create GitHub accounts with guidance from Bryan (2018b) on choosing a user name. At the beginning of the semester, a portion of lecture is used to introduce version control and reproducibility, why they are important, and how RStudio and Git will help students implement these practices in their work.

One of the first assignments in the course is a computing assignment focused on using RStudio and Git. This assignment serves as a review for some students, and it is an introduction to these tools for others. For all students, however, it is their first exposure to the workflow they will use throughout the rest of the course. Students write their responses in an R Markdown file, knit the file to produce a Markdown document, write a short and informative commit message, and push their work to GitHub for submission. Throughout the assignment instructions are periodic reminders to *knit*, *commit*, and *push*, a mantra used throughout the semester to remind students how to connect their work in RStudio to their assignment repository in GitHub. The instructions for the first few assignments also include examples of informative commit messages. The instructions end with a reminder for students to review their work in the assignment

repository on GitHub to ensure it is the final version to be submitted for grading.

Students complete several assignments individually before working on their first team assignment. This gives them an opportunity to become familiar with RStudio and Git and become comfortable with this workflow before introducing the additional layer of collaborating in GitHub. For the first team assignment, in addition to the aforementioned workflow cues, students also receive cues to *pull* so they have the most updated version of the collaborative document. There are also cues to rotate which team member types the responses. These workflow cues are eventually removed from the assignment instructions as the semester progresses and the workflow is more routine for students.

5.4. Workflow

There are two basic workflows in the course: one for general assignments, such as homework and computing labs, and one for the final project.

The typical workflow for general assignments is the following:

1. The instructor creates a starter repository. The starter repository includes a link to the assignment instructions, an R Markdown template, and a folder for the data (if needed). In the beginning of the semester, the dataset is already included in the starter repository. As the semester progresses, students download the data from the assignment instructions and upload it to the repository.
2. A copy of the starter repository is created for each student (or team) using the `ghclass` R package. For individual assignments, the repositories are named using the template `assignment_name-[user_name]`, where `[user_name]` is the student's GitHub username. For team assignments, the repositories are named `assignment_name-[team_name]`, where `[team_name]` is the team's name on GitHub. For example, the first individual homework assignment is named `hw01-[user_name]`.
3. Students start a new project in RStudio by cloning their assignment repository. They configure the RStudio project with the GitHub repository by using the `use_git_config()` function in the `usethis` R package (Wickham and Bryan 2020). Students complete the assignments in RStudio, by typing their responses in an R Markdown document with output: `pdf_document` which produces a PDF from the R Markdown document. They periodically *knit*, *commit*, and *push* their work to their repository on GitHub.
4. Students submit their work by connecting their repository to the associated assignment on Gradescope, an online rubric and grading system (Gradescope 2020).
5. Students view the assignment feedback on Gradescope. It is connected to the LMS which ensures that grades are securely stored within the university's system.

During the second half of the semester, students complete a final project in teams of three or four. The workflow for the project is generally similar to the one described above, with the exception being how students receive feedback. At various checkpoints in the project, students receive feedback as an

"issue" in the GitHub repository for their project. They can reply to the issue as one way to ask the instructor follow-up questions about the comments. This feedback workflow is used for the project to more closely mimic how students may exchange ideas with collaborators if they use GitHub outside of the classroom setting. Though comments are posted in the GitHub issue, there are no grades posted on GitHub. All grades related to the project are posted only in the LMS.

5.5. Assessment

Developing a proficiency using RStudio and Git is a learning objective for the course, so students are assessed on how they use the tools on a majority of assignments. Students are required to have their work in their GitHub repository to be considered for grading, so they must learn how to *push* to GitHub to complete individual assignments and both *push* and *pull* for team assignments.

Each assignment includes a category named "Overall" that includes points dedicated to using Git. Typically about 5% of the points on an assignment are for having at least three commits and writing informative commit messages. On team assignments there are also points allocated for having at least one commit from each team member. This is used to hold team members accountable for contributing and to encourage teams to make use of the collaborative nature of GitHub.

5.6. Other Remarks

Based on multiple semesters of teaching version control in undergraduate courses, here are a few recommendations for instructors who are considering teaching GitHub as a learning objective in an undergraduate statistics course:

- Get early buy-in from students. As mentioned in Section 5.3, a portion of lecture in the beginning of the semester introduces the importance of reproducibility and version control. Given the relatively steep learning curve for Git and GitHub, it is important that students understand the value of learning these computing skills and how they are used when doing a statistical analysis.
- Focus only on the GitHub functionality used in the course, as introducing too much functionality can become overwhelming. Generally teaching students how to *push*, *pull*, *commit*, and resolve merge conflicts is enough to complete the assignments in a statistics course.
- Use the Git pane in RStudio. Running Git commands through the RStudio interface helps make it more accessible for students who don't have previous experience running code from the command line.

6. A Master's Level Statistical Programming Course (Rundel)

6.1. Course Description

STA 523: Statistical Computing was developed in 2016 for the (then) new Master's in Statistical Science (MSS) program at Duke University and has been offered yearly since. The course

was designed around three core pillars: focusing on reproducible methods, emphasizing programming knowledge, and teaching foundational data science skills. The course shares many commonalities with approaches to integrate computing suggested by Nolan and Temple Lang (2010) (see <https://www.stat.berkeley.edu/~statcur/>).

The course is required for all first year MSS students in their first semester, and consists of two 75-min lectures and one 75-min workshop per week. These students come from a wide variety of backgrounds with many having little or no prior coding experience. Similarly, most students have never used or have had minimal experience with Git and GitHub or other version control systems. As the only required course in the MSS program that focuses on computing, programming, and software engineering, the goal of the course is to provide the students with a strong foundation of skills that are relevant to their other courses as well as their careers after graduation. While the ideal computational skill set for an MS in Statistics graduate is a moving target, it has become clear that beyond traditional topics (e.g., numerical computing, optimization, etc.) more data-focused skills (e.g., data munging, databases, SQL, etc.) are increasingly important. We have attempted to reflect this in the course's evolving curriculum. In addition, the course covers statistical topics such as modeling and prediction as well as Bayesian methods such as approximate Bayesian computation and MCMC. These statistical topics are presented to complement other coursework in the curriculum by focusing on the computational details and implementation.

6.2. Tools and Implementation

Similar to the two previous courses, students in this course use the RStudio IDE and interact with GitHub via the Git pane in RStudio. In the earliest iterations of the course the process for creating, distributing, and collecting student work from GitHub repositories was done manually or via simple shell scripts. Over time a number of tools have made this process much easier by allowing for the automation of most of these processes. Specifically, GitHub released their Classroom tool around the same time this course was first offered and it was used for the delivery of individual assignments. However, its team assignment workflow was initially not present and later overly constraining, as it did not allow the instructor to assign teams.

These limitations and the power and availability of the GitHub API led to the development of the `ghclass` R package which is used for automating interactions with GitHub for course management (Rundel, Çetinkaya-Rundel, and Anders 2020). The package has more functionality than can be explored in this article, but the core use case is for automating the creation of team and individual assignment repositories. For example, using the template repository based workflow, described above, the package helps create new repositories, create teams, add members, add teams or users to the repositories, and then copy the template's contents with a single function, for example, `org_create_assignment()`.

One additional attractive feature of the package is the ability to interact with existing repositories, particularly when it comes

to adding or modifying files. This is valuable to address the all-too-common situation that a distributed assignment includes a typo, a minor issue with the data, or the repo contains the wrong version of a file. Rather than having to send out an E-mail announcing the issue or posting an announcement to the course LMS, `ghclass` allows for *pushing* the corrected file(s) out to all of the students' repositories in a way that is merged with any existing work. Since everything is managed via Git there is no risk of, permanently, overwriting student work. The work that has gone into implementing the necessary low level functionality in this package has allowed extensions and experiments with the automation of higher level processes (such code formatting feedback and peer review).

As some of the topics in the course involve fairly heavy computational workloads, a centralized powerful departmental server hosting RStudio Server Pro provides students access to the necessary compute resources. The more extensive computational workload is also relevant to the use of large datasets in the course, as GitHub has a strict file size limit of 100 MB, while some of the datasets used in this class exceed multiple gigabytes in size. To address this issue, these large files are hosted in a shared, read-only directory on the same server that hosts RStudio, such that all students have access to the files without having to maintain their own copy. These data could also be hosted in the cloud, but co-locating the data with the compute resources is important both for efficiency and cost.

6.3. First Exposure in Class

As stated previously, the expectation is that students will use Git and GitHub for all of their assignments within the class starting from day one. To ease the initial learning curve for these tools, as well as RStudio, the class explicitly includes at least 1 hr of lecture in the first week to motivate these topics. Typically, this takes the form of a very brief introduction to the theory and history of version control and Git, and then the remainder of the time is dedicated to a live demonstration of the tools.

At this time, the first assignment is distributed via a GitHub Classroom link: students follow the link, connect their GitHub account to a unique identifier in the roster, and then gain access to their own private repository copy of the assignment template repository. As part of the live demonstration, the instructor leads students through this process and how to locate and interact with the repository on GitHub. This leads to forking the repository, cloning the Git repository as an RStudio project, and a demo of the basic use of the Git pane. The basic Git actions such as *stage*, *commit*, *push*, and *pull* are covered and the class usually concludes by purposefully inducing a merge conflict to demonstrate the process of resolving it.

This is a large amount of material for the students to absorb in a short period of time. Recording the session (either via screen recording or lecture capture), providing in-person support, and having an initial individual assignment that is focused on reinforcing the workflow has allowed most students to get up to speed quickly. The number of Git-related questions during workshops and in office hours is substantial during the initial weeks of each semester, but tends to decrease rapidly as the semester progresses.

6.4. Workflow

The course features four to eight team assignments and two individual projects, all of which require complete reproducibility for full credit. For each assignment/project a template repository is created, which is structured using an RStudio project and contains the files necessary for the assignment. Typically, this template includes a README file that contains a detailed description of the assignment, a scaffolded R Markdown document that gives a uniform structure for the assignment and includes clear indications of where the students should enter their implementations/solutions and write up, and any additional necessary support files (e.g., data, scripts, etc.). These template repositories can easily be shared with teaching assistants for feedback and or training purposes, and can also be moved from previous years into new organizations for each new offering of a course. The work is assigned to the students by mirroring the template repository to either the student or their team's private repository, using a consistent naming scheme, for example, `hw01-team01`, which gives them access to their own copy of all the necessary files for the assignment and can be directly cloned as an RStudio project from GitHub using the New Project interface.

Students are then able to work on the assignment within RStudio and turn in their work as well as collaborate with team members by *committing* and *pushing* their code back to GitHub. Earlier versions of the course focused on teaching both the RStudio Git interface as well as the command line Git interface, but the vast majority of students preferred the former so the latter approach was dropped in favor of adding other content. At the deadline for each assignment it is simply a matter of cloning all of the assignment repositories from the organization to obtain a local copy of all of the students' work, which can then be rerun to assess reproducibility and the resulting HTML or PDF documents graded.

6.5. Assessment

While grading each student or team's work, their R Markdown documents are recompiled to ensure the reproducibility of their work. In early versions of the course this was coupled with a course policy that work that failed to compile would receive a zero, which turned out to be almost impossible to enforce in practice. Most of the time students' code would fail to run for relatively small issues (e.g., use of `setwd()` to set working directories with absolute file paths or loading a less commonly used package that the instructor did not have installed) that could be easily fixed yet caused a compilation error. It was possible to have a back and forth with the students about the errors and have them correct them but this proved to be very inefficient and frustrating for both the students and instructor.

The solution to this has been to implement automatic feedback for the students on the basic processes of their assignment. This is done by checking their R Markdown documents and repositories using continuous integration tools available via GitHub Actions. These tests take the student's code and run basic sanity checks every time students *push* their code to GitHub: does the R Markdown document knit, are only the

necessary files included in the repository? GitHub Actions are used to implement these tests. The test results are signaled to students via a badge in their repository README that shows either green or red depending on whether the check passed or not. Additionally, they can click on these badges to get specific feedback in the case a check failed. This becomes a simple necessary (though not sufficient) condition for the students to examine when completing an assignment. Examples of the current GitHub action based workflows being used for this and related courses are available at the `ghclass-actions` GitHub repository (Rundel 2020).

Use of Git and GitHub has never been an explicitly assessed component of these courses, however, it is inherently tied to the students work as it is the only available method of obtaining and turning in their work. There have not been any specific efforts to encourage particular workflows with Git/GitHub (i.e., branching, issues, etc.) but it has been interesting to observe some of the emergent behaviors that student teams have developed for collaboration.

6.6. Other Remarks

This specific course has been offered every Fall since 2016 and has consistently had a capped enrollment of around 40 students per semester. In 2017, it was decided to add an undergraduate equivalent to the course, STA 323, which has been offered in the Spring semester each year since and has also had a capped enrollment of approximately 40 per semester. Lastly, upon joining the University of Edinburgh, a similar Master's-level Statistical Programming course, Math 11176, has been taught which has an enrollment of around 180 students.

The example of Jenny Bryan and her Stat 545 course (Bryan 2020) at the University of British Columbia served as a direct inspiration for these courses. It was invaluable to see that such a course already existed and had been successful over a number of years. The open publishing and dissemination of the course materials directly influenced many aspects of the class. Without this clear model of a course it is unlikely that these courses would have been developed or been as successful. More of this type of creative sharing would be of benefit to the community.

7. A Master's Level Biostatistics Course (Sullivan)

7.1. Course Description

Statistical Programming in R (PHP 2560) is a first programming course for Master's students in Biostatistics at Brown University. The main focus of the course is to develop good statistical programming habits while also preparing students for the world of reproducible research and data science. The class started with 14 students and as of Fall 2019 drew more than 50 students a semester, including students from undergraduate to Ph.D. level in multiple departments across the University. Many students come to the course with some experience using R for data analysis, but no experience writing functions, designing packages, or creating interactive web applications with the `shiny` package (Chang et al. 2019).

7.2. Tools and Implementation

Students use their own computers and have the choice between installing both R and RStudio locally or via RStudio Cloud. Each student creates a GitHub account and connects this to their RStudio use. Pre-class and in-class coding exercises are all placed into GitHub repositories that act as starter code for assignments created using GitHub Classroom. When students accept the assignment, they use RStudio to clone their GitHub repository as an RStudio project. They are guided to *commit-pull-push* frequently as they complete their work. As noted earlier, GitHub Classroom provides a flexible platform for creating private GitHub repositories for students.

7.3. First Exposure in Class

The course is about statistical programming in R, however, the first few class meetings are dedicated to learning Git. Prior to any instruction in R, the students work through the basics of Git. Before the first class, students install Git, R, and RStudio on their computers, then work through the “First days of Git” Learning Lab (GitHub 2020b). During the first class, students follow a link to a starter assignment from the course website and they create an RStudio project with this link. Then, students work through a few exercises throughout the entire first class practicing the Git workflow and RStudio.

7.4. Workflow

The rest of the course consists of DataCamp (DataCamp 2019) assignments and pre-class coding exercises. The pre-class work is all in one repository to which students contribute throughout the semester. The course website contains links to either GitHub repositories for students to grab files from or an assignment link generated using GitHub Classroom. The class consists of a 15 minute overview and review of pre-class work. Then students begin a project by cloning their assigned project in RStudio. Each class they first *commit* their pre-class work into the team repository. They then review and comment on each others code and *push* those comments and feedback. At this point, they begin working on team coding projects. They typically create their own scrap work file and then together decided on which code is displayed on the teams’ final results.

When they encounter their first merge conflict, they are instructed to work through this with the help of the instructor as well as GitHub’s merge conflict tutorial (GitHub 2018). The Git process is slow going and some students find it to be frustrating at first, but it does not take long before they have very few conflicts or problems with their repositories.

7.5. Assessment

The instructor and teaching assistants create plain text files (e.g., `feedback.md`) in each repository and comment the code based on a published rubric that all students see on the course site. Scoring rubrics include evaluation of the repository history and commit messages as well as R programming style. Feedback

is then added to the student repositories so they can *pull* to see the remarks shared with them.

7.6. Other Remarks

The experience of teaching this course for over 4 years (6 offerings of the course) led to the following recommendations for other instructors considering teaching a course with version control as a learning objective:

- Experiment with Git and RStudio implemented on various computers. If students run into problems while installing or configuring the tools, it is helpful to have experience with more than one implementation. Use students who have things working on their computer to help others troubleshoot. This creates teamwork but also allows for more students to receive support at a time.
- Invest some time early to motivate the workflow with Git and address common pitfalls. For example, students should avoid committing files that they did not actually change to help minimize merge conflicts, and instead investigate why they show up in the Git pane in the first place. Each student working in shared repositories could be encouraged to maintain an individual scrap work file in the repository that only they edit as a measure to avoid merge conflicts or overwriting someone else’s work. Such files should be removed from the repository before final submission.
- Be patient. Students and instructors alike may encounter challenges in the beginning, and sometimes it can be hard to diagnose the problem they are having. However, you will find that their ability to code in teams and track their work begins to outweigh any issues.

8. Discussion (Çetinkaya-Rundel and Horton)

In addition to the organizing prompts, each contributor populated a matrix of learning outcomes with codes representing the type of exposure typical in the course(s) described. Table 2 presents this matrix using the following symbolic representation for each learning outcome in each course:

- □□□□: None. This is not included in course.
- ■□□□: Incidental. This may or may not occur in the course.
- ■■□□: Teacher. This is demonstrated to students, but they are not necessarily expected to do it independently.
- ■■■□: Student. Students are expected to do this independently, but it may not be formally assessed.
- ■■■■: Assessed. Students are expected to do this independently, AND will be assessed for proficiency.

The instructors of the undergraduate courses had similar approaches to the use of GitHub, with more advanced topics (e.g., branching and continuous integration) only showing up in the graduate courses. There was considerable heterogeneity between topics in terms of whether they were formally assessed.

While there are a number of commonalities to these instructor stories, there are also some differences. In this section we provide a high level overview of some of the issues raised.

Table 2. Git learning outcomes and assessment across the courses.

Learning outcome	Beckman	Tackett	Rundel	Sullivan
Repositories:				
clone a private repository and push a commit	■ ■ ■ ■	■ ■ ■ ■	■ ■ ■ ■	■ ■ ■ ■
create a repo	■ ■ ■ ■	■ ■ ■ ■	■ ■ ■ ■	■ ■ ■ ■
create a branch, merge branches	□ □ □ □	□ □ □ □	■ ■ ■ ■	■ ■ ■ ■
retrieve an older version of a file	■ ■ ■ ■	■ ■ ■ ■	■ ■ ■ ■	■ ■ ■ ■
continuous integration or other automation	□ □ □ □	□ □ □ □	■ ■ ■ ■	□ □ □ □
GitHub Issues:				
create, comment on, and/or assign an issue	■ ■ ■ ■	■ ■ ■ ■	■ ■ ■ ■	■ ■ ■ ■
reference commits/code line numbers in issues	□ □ □ □	■ ■ ■ ■	■ ■ ■ ■	■ ■ ■ ■
Collaboration:				
student teams collaborate in a shared repo	■ ■ ■ ■	■ ■ ■ ■	■ ■ ■ ■	■ ■ ■ ■
resolve a merge conflict	■ ■ ■ ■	■ ■ ■ ■	■ ■ ■ ■	■ ■ ■ ■
fork and create a pull request	■ ■ ■ ■	□ □ □ □	■ ■ ■ ■	■ ■ ■ ■
merge a pull request	■ ■ ■ ■	□ □ □ □	■ ■ ■ ■	■ ■ ■ ■
review changes and blame	■ ■ ■ ■	■ ■ ■ ■	■ ■ ■ ■	■ ■ ■ ■
create gh-pages	■ ■ ■ ■	□ □ □ □	■ ■ ■ ■	■ ■ ■ ■

8.1. Students Need to See Value of These Expert-Friendly Tools

As instructors it can sometimes be frustrating to teach foundational tools and approaches since students often want to jump directly to fancy models or visualization. This may leave them unable to carry out simpler and more straightforward tasks where their analyses can be documented and reviewed.

How can we help to motivate students to think about the importance of workflow and develop internal motivation? Peter Norvig (Google) notes that what students need most is “meticulous attention to detail” (National Academies of Science, Engineering, and Medicine 2019). Are there ways that we can help them develop and strengthen this capacity by demonstrating that source code control is a tool that can be useful to tracking their work and to help them be less error-prone? One approach might be to share a cautionary tale, perhaps Xiao-Li Meng’s story (Meng 2020) of the data loss of much of his doctoral dissertation.

We saw multiple examples of instructor scaffolding to provide a guided introduction to the power and value of GitHub (to track and document their work) without getting lost in the details. We believe that the scaffolded introduction to GitHub is a useful if not sufficient framework to build more habits that foster reproducibility.

8.2. Start Slowly and Keep It Simple

A key takeaway of the approaches described in the article is how instructors have started slowly and gradually built up complexity. The instructors adopted a “less is more” approach to avoid cognitive overload. This is evidenced in how they each structure students’ first exposure to version control as well as how they (almost) all limit interactions with version control to a small number of Git actions through the RStudio IDE. While courses that feature teamwork get into thorny concepts like merge conflicts, these are deferred until later in the course, after students develop more comfort managing version control.

Other more advanced features of Git (e.g., branching, pull requests, rebasing, HEAD) are both valuable and commonly used by data scientists, but such details might be appropriate to leave until later. Having students learn to use straightforward Git workflows early on is a big step on their path to developing good habits for workflow and collaboration.

We should also note that keeping it simple does not necessarily mean that students will not learn about more advanced Git tricks from other resources. For example, it is possible to backdate a Git commit. Since there is no notion of a “deadline” in Git repositories, students could presumably backdate a commit made after the deadline for an assignment as though it was made before the deadline. It is, however, possible to prevent students from making changes to their repositories after a deadline via a few indirect methods. Instructors can collect (clone or download the contents of) student repositories at the deadline. Alternatively, instructors can change permission levels of students at the deadline so that they can no longer *push* changes, but can continue to read and interact with the repository. Both of these methods can be automated using the `ghclass` package.

8.3. Why GitHub (and Not GitLab, Bitbucket, etc.)?

There are a number of web-hosting platforms for projects version controlled with Git. The three most popular of these platforms are GitHub, GitLab, and Bitbucket. Among these, GitHub is recognized as the industry standard platform for hosting and collaborating on version controlled files via Git with an estimated more than 2.1 million businesses and organizations using GitHub (GitHub 2020a), compared to an estimated number of one million users on Bitbucket (Bitbucket 2020) and more than 100,000 organizations on GitLab (GitHub 2020a).

In addition, GitHub provides a rich API, which allows for tools like GitHub Classroom and `ghclass`. The `ghclass` package offers functionality for peer review by moving files around between GitHub repositories of students. Additionally, features like GitHub Actions can be used for immediate feedback and auto-grading of files by triggering certain code to run in the background every time students *push* to their repositories.

One potential difficulty with using GitHub is the fact that the student code—even in private repositories—is hosted on GitHub servers, which means student data leaves the university. This is especially important for institutions outside of the United States since there may be laws around student data leaving the country and being stored on US servers, for example, the European Union’s General Data Protection Regulation (GDPR). One remedy to this is for the university to enter a data protection agreement with GitHub for GDPR compliance. Another

possibility is for the university to host their own GitHub server. The software associated with this, GitHub Enterprise, is freely available for academic teaching use. However, the university needs to supply the hardware (server) as well as IT resources to set up the server and student authentication. Note that the latter is a much more resource intensive solution.

We note an important potential danger of building curricula solely around any specific commercial technology, including GitHub. Until 2018, GitHub was a start-up. That year it was acquired by Microsoft. It's impossible to tell what is next for the company. The company currently seems dedicated to continue offering free private organizations and repositories for educational use, and we have no reason to doubt this. However, it serves as a reminder that software tools and their terms of use do change from time to time. It is certainly a risk, but we note that if you are teaching data science, and want to stay current, you should be willing to take a certain amount of risk, in a calculated way such that the students do not end up suffering consequences harshly. Many educators and developers are investing time in building infrastructure and tooling to help others stay current with their data science pedagogy and tooling. Instructors, even those not interested in participating in the development of such tools, should track what is being developed to ensure their programs stay current in this rapidly changing environment.

8.4. Not One Single Path

One striking take-home observation from the instructors' stories is that different instructional teams follow different models with different pedagogical goals, ranging from fostering collaboration or automating aspects of course mechanics. Some of the ways students and instructors use GitHub include: (1) a pull request model, (2) full write access to individual or team project repos, (3) use of the `ghclass` package with one repository per student per assignment, or (4) use of GitHub Classroom with one repository per student per assignment.

There are also various approaches to assessing student work and providing feedback on GitHub. Use of issues to provide feedback is a popular approach, and it is possible to make this process more efficient and streamlined with the use of issue templates. Automated feedback using continuous integration tools like GitHub Actions is another approach that can either replace or supplement the manual feedback process via issues. Writing automated tests to fully evaluate data science assessments that not only contain code and output but also interpretations is difficult, and perhaps impossible. Therefore, it is difficult to imagine how automated checks can fully replace manual grading, but developments on this front are exciting to see as statistics and data science classes grow in size.

A key implication of these various approaches is that we do not want to be too prescriptive in terms of a specific workflow. We might consider an analogy to code style: there are multiple reasonable ones, and we can (and do) engage in somewhat religious arguments about what is "right," but the key aspect is that there are compelling reasons to "fit in." The same approach is important when thinking about teaching GitHub and version control.

Some of the structures described in the article (e.g., the `ghclass` package), are powerful and flexible systems that facilitate scaling to larger classes. As a reviewer notes, these systems have a nontrivial learning curve. For classes with no team-based work, and especially for an instructor who is just starting with Git and GitHub, we recommend using GitHub Classroom for managing the distribution and collection of students assignments as repositories. For classes that also involve teamwork, the `ghclass` package offers more complete functionality for course management. Additionally, the `ghclass` package also provides support for automation of editing and correcting repositories that have already been distributed to students as well as automating many other common tasks that need to be applied across a large number of repositories via the GitHub API, for example, managing organization and team membership, retrieving repository statistics, testing using GitHub Actions, etc. As of the time of writing this article, these features are not offered in GitHub Classroom. It should also be noted that GitHub Classroom and the `ghclass` package are compatible tools and can be used in alongside each other.

8.5. Peer Review

Another approach that several of us are exploring in our courses is peer review, which has the benefit of exposing students to each others' work and also prepares them for industry settings where code review is commonplace. GitHub is already designed for peer review as this is a crucial part of software development. From a technical perspective, peer review is enabled with the functions starting with the `peer_*` prefix in the `ghclass` package. They offer functionality for retrieving files from one repository, anonymizing by stripping the metadata (e.g., student names, commit history), moving these files to a new repository where a randomly selected student has access to read and review, and collecting these reviews and submitting them as a pull request to the original student repositories. The full peer review functionality and process is described in detail in a vignette in the `ghclass` package (Rundel, Çetinkaya-Rundel, and Anders 2020). Peer review gives students the opportunity to meaningfully engage with each others' work and learn from each other. It also gives them a chance to try to reproduce others' work and experience the difficulties of reproducing others' work.

But the workflow that is native to GitHub (via branches, pull requests, and no anonymity) does not always work for teaching—either because some of these concepts are beyond the learning goals of the course (e.g., in intro courses we do not talk about branches and pull requests) or may not be suitable for a learning environment (an instructor might want to do anonymous peer review so students do not know whose work they are reviewing).

GitHub's rich API allows us to leverage what's already built in to GitHub and customize it to be more suitable for peer review as part of coursework. Several of the authors are exploring how best to integrate peer review facilitated by GitHub into our courses.

8.6. Creating Portfolios

Students sometimes use their GitHub profile as part of their job search (Tech Beacon 2020). Educators may encourage their

students to curate their GitHub profile based on their coursework. Use of GitHub in courses may assist with this process. However, this is not automatic as coursework needs to be stored in private repositories and it is not always obvious how to expose this work while conforming to FERPA, GDPR, etc.

One approach is to only allow students to convert their final project repository to be public, and let them know of specific assignments from class that are approved for reproducing in public repositories. These are usually assignments with low risk of plagiarism and/or high personalization. Moreover, it is important for instructors to keep summative or overly critical feedback and any grades out of repositories that students might convert to public repositories later. Any team work also needs to be handled with care: all team members need to agree that work can be made public.

It should also be noted that it can take time for students to get their class projects into something that is portfolio worthy. Often times, a repository with just some code does not make a compelling portfolio entry. Students need to create an informative but brief write-up that features highlights from their work, so that those browsing their portfolio know where to start looking, or more importantly, why this repository is worth looking into. One quick solution for this is a rich README. An additional step is to publish the repository as a webpage simply by turning on GitHub Pages (gh-pages) feature, which will turn the README of a repository into a webpage. This process can be an official part of an assignment or provided to students as a parting gift to help increase their visibility on the web. Future research in this area might explore ways that such e-portfolios might be helpful in both curating student work and highlighting their efforts.

8.7. Assessment

Many of the instructors have touched on the commit history providing a transparent account of the work done by students in their repositories, which can be especially useful for individual accountability in teamwork. While number of commits, on its own, is not a strong indicator of the quality of work done by a student, lack of commits can signal that the student has not made direct contributions to a team project. However, unless the course makes it clear that commits by each team member are required, this number alone might not be a true representation of a student's contribution. For example, if pair programming without switching roles, commits would appear to be made only by one student. We recommend making use of peer evaluations in any courses that involve teamwork to get a clearer picture of each student's contribution, and supplementing the feedback from these evaluations with commit history statistics.

Since version control is a learning goal for these courses, we believe that instructors should assess how well students are following recommended workflows in each assignment. We recommend assigning roughly 10% of the points in each assignment to organization and style (e.g., figure sizing, code style, formatting, etc.), and a portion of those points specifically to version control related tasks. These include: (1) reasonable number of commits, (2) reasonable commit content, and (3) meaningful commit messages. Getting statistics like number of commits is

made possible with the `ghclass` package. Assessing reasonable commit content can be a lot more cumbersome, and likely only worth looking into if the student has too few commits on an assignment. Finally, for assessing whether the commit messages are meaningful, we recommend quickly taking a peek at the commit history on the GitHub repository, and scanning to see if there are any commit messages that don't obviously meet this criterion (e.g., random string of characters or too many commits that just say "update").

It is crucial for instructors to model good version control hygiene before assessing it, as students cannot be expected to come into the class with an intuition for what is reasonable commit content or message. One way of doing this is explicitly stating when and what to commit in earlier assignments (e.g., "make a commit after this exercise") and what to say in the commit message. Throughout the semester this sort of scaffolding can be slowly removed from assignments, letting the students learn to make a decisions about what constitutes a reasonable change that should be captured in a commit. (It is also imperative that the instructor practice what they preach in terms of instructor commits in course and student repositories.)

8.8. Automation and Workflow

More advanced users (and many instructors) will benefit from the use of automation tools. For the instructor, this might facilitate auto-pulling local files, simplify returning feedback to students (e.g., a script to open all repositories to the issues page, a script to pull files, add a file called `feedback.md` that includes comments, `commit`, and `push` to many repos). The use of continuous integration tools to check for compilation may be particularly helpful as students work on more complex tools and approaches that go beyond the capabilities of their laptops. This is an area where new approaches are being developed in the research community that have the potential to improve student experiences and/or simplify work by the instructor and improve learning outcomes for students.

We hope that this article encourages instructors to explore and share their experiences.

8.9. Closing Thoughts

Students heading into the workforce need to be able to structure, organize, and communicate their work. Using version control is a valuable, useful, and now logistically practical tool, so we recommend instructors consider incorporating it into their courses and programs.

It is worth mentioning that there are plenty of viable alternatives to the RStudio IDE (e.g., Atom, JupyterHub, Vim) or Git for version control (e.g., Subversion, Mercurial). Several of these tools have similarly efficient integration, and it should be clear that no attempt was made to compare and contrast alternative implementations in this article.

One limitation of this article is that it only provides the educator's point of view to using Git and GitHub. We have observed a few patterns emerging in how students work with version control and how they collaborate over a version control platform. Students tend to figure out the basics of working with version

control on individual assignments pretty quickly, within the span of a few weeks. However, they often find collaboration, and especially merge conflicts, more challenging. To help ease this challenge students may get together in person for team projects so that they can make commits from a single computer, which can also be seen as a positive outcome, but it shows that students find collaborating on GitHub challenging. Additionally, whether there is long term adoption of using version control by the students is less clear. However, we observe that students coming out of these courses and then working on research projects or participating in ASA DataFest, a weekend-long, team-based data analysis competition (Gould and Çetinkaya-Rundel 2014), often choose to use version control and collaboration with Git and GitHub.

Finally, version control has not come up as an issue students bring up in course evaluations; in fact, it regularly gets mentioned positively among skills they learned in the courses. None of the classes mentioned in this article have systematically collected data on student attitudes toward version control. We believe that this would be a valuable next step for statistics and data science courses so that we can explore how the implementation of GitHub in the classroom is associated with students' classroom experiences, similar to Hsing and Gennarelli (2019), which discusses such a study conducted on students in computer science courses. Such assessment data and other findings informed by the learning sciences would help improve instruction in this area. Future research could help inform a publication akin to Hesterberg's (2015) guide to teaching resampling entitled "What every statistics and data science instructor should know about version control and reproducible workflows."

ORCID

Mine Çetinkaya-Rundel  <http://orcid.org/0000-0001-6452-2420>

Nicholas Horton  <http://orcid.org/0000-0003-3332-4311>

References

- American Statistical Association (2014), "Curriculum Guidelines for Undergraduate Programs in Statistical Science," available at <http://www.amstat.org/education/curriculumguidelines.cfm>.
- Baumer, B., Çetinkaya-Rundel, M., Bray, A., Loi, L., and Horton, N. J. (2014), "R Markdown: Integrating a Reproducible Analysis Tool Into Introductory Statistics," *Technology Innovations in Statistics Education*, 8, available at <https://escholarship.org/uc/item/90b2f5xh>.
- Bitbucket (2020), "Bitbucket," available at <https://bitbucket.org>.
- Bryan, J. (2018a), "Excuse Me, Do You Have a Moment to Talk About Version Control?," *The American Statistician*, 72, 20–27.
- (2018b), "Happy Git and GitHub for the useR," GitHub, available at <https://happygitwithr.com>.
- (2020), "Stat 545 Website," available at <https://stat545.com>.
- Çetinkaya-Rundel, M. (2020), "Data Science in a Box," available at <https://www.datasciencebox.org>.
- Çetinkaya-Rundel, M., and Rundel, C. (2018), "Infrastructure and Tools for Teaching Computing Throughout the Statistical Curriculum," *The American Statistician*, 72, 58–65, DOI: 10.1080/00031305.2017.1397549.
- Chang, W., Cheng, J., Allaire, J., Xie, Y., and McPherson, J. (2019), "shiny: Web Application Framework for R," R Package Version 1.4.0, available at <https://CRAN.R-project.org/package=shiny>.
- DataCamp (2019), "DataCamp," available at <https://datacamp.com>.
- Federal Educational Rights and Privacy Act (FERPA) (2020), "Federal Educational Rights and Privacy Act (FERPA)," available at <https://www2.ed.gov/policy/gen/guid/fpco/ferpa/index.html>.

- Fiksel, J., Jager, L. R., Hardin, J. S., and Taub, M. A. (2019), "Using GitHub Classroom to Teach Statistics," *Journal of Statistics Education*, 27, 100–119.
- Garfield, J., Zieffler, A., Kaplan, D., Cobb, G. W., Chance, B. L., and Holcomb, J. P. (2011), "Rethinking Assessment of Student Learning in Statistics Course," *The American Statistician*, 65, 1–10.
- GitHub (2018), "GitHub Learning Lab," available at <https://lab.github.com/githubtraining/managing-merge-conflicts>.
- (2020a), "GitHub," available at <https://github.com>.
- (2020b), "GitHub Learning Lab," available at <https://lab.github.com>.
- GitHub Education (2020), "GitHub Classroom," available at <https://classroom.github.com>.
- Gould, R., and Çetinkaya-Rundel, M. (2014), "Teaching Statistical Thinking in the Data Deluge," in *Mit Werkzeugen Mathematik und Stochastik lernen—Using Tools for Learning Mathematics and Statistics*, eds. T. Wassong, D. Frischmeier, P. Fischer, R. Hochmuth, and P. Bender, Wiesbaden: Springer Spektrum, pp. 377–391.
- Gradescope (2020), "Gradescope," available at <https://www.gradescope.com>.
- Haaranen, L., and Lehtinen, T. (2015), "Teaching Git on the Side: Version Control System as a Course Platform," in *Proceedings of the 2015 ACM Conference on Innovation and Technology in Computer Science Education, ITiCSE '15*, ACM, New York, NY, USA, pp. 87–92.
- Hardin, J. S., Hoerl, R., Horton, N. J., Nolan, D., Baumer, B., Hall-Holt, O., Murrell, P., Peng, R. D., Roback, P., Lang, D. T., and Ward, M. D. (2015), "Data Science in Statistics Curricula: Preparing Students to 'Think With Data,'" *The American Statistician*, 69, 343–353.
- Hesterberg, T. (2015), "What Teachers Should Know About the Bootstrap: Resampling in the Undergraduate Statistics Curriculum," *The American Statistician*, 69, 371–386.
- Hsing, C., and Gennarelli, V. (2019), "Using GitHub in the Classroom Predicts Student Learning Outcomes and Classroom Experiences: Findings From a Survey of Students and Teachers," in *Proceedings of the 50th ACM Technical Symposium on Computer Science Education, SIGCSE '19*, ACM, New York, NY, USA, pp. 672–678.
- Kaggle (2017), "Kaggle Machine Learning & Data Science Survey 2017," available at <https://www.kaggle.com/kaggle/kaggle-survey-2017>.
- Kaplan, D. T. (2015), "Data Computing: An Introduction to Wrangling and Visualization With R," Project Mosaic.
- Kaplan, D. T., and Beckman, M. D. (2019), *Data Computing* (2nd ed.), available at <https://dtkaplan.github.io/DataComputingEbook>.
- Meng, X.-L. (2020), "XL-Files: Time Travel and Dark Data," *IMS Bulletin*, 49, 6.
- National Academies of Science, Engineering, and Medicine (2018), "Data Science for Undergraduates: Opportunities and Options," available at <https://nas.edu/envisioningds>.
- (2019), "Roundtable on Data Science Postsecondary Education Meeting 10," available at <https://nas.edu/dsert>.
- Nolan, D., and Temple Lang, D. (2010), "Computing in the Statistics Curriculum," *The American Statistician*, 64, 97–107.
- RStudio Team (2015), *RStudio: Integrated Development Environment for R*, Boston, MA: RStudio, PBC, available at <http://www.rstudio.com>.
- (2020), *RStudio Cloud*, Boston, MA: RStudio, PBC, available at <https://rstudio.cloud>.
- Rundel, C. (2020), "ghclass Actions," available at <https://github.com/rundel/ghclass-actions>.
- Rundel, C., Çetinkaya-Rundel, M., and Anders, T. (2020), "ghclass: Tools for Managing Classes With GitHub," available at <http://github.com/rundel/ghclass>.
- Tech Beacon (2020), "What do Job-Seeking Developers Need in Their GitHub?," available at <https://techbeacon.com/app-dev-testing/what-do-job-seeking-developers-need-their-github>.
- Wickham, H., and Bryan, J. (2020), "usethis: Automate Package and Project Setup," available at <https://github.com/r-lib/usethis>.
- Xie, Y., Allaire, J., and Golemund, G. (2018), *R Markdown: The Definitive Guide*, Boca Raton, FL: Chapman and Hall/CRC.
- Zagalsky, A., Feliciano, J., Storey, M.-A., Zhao, Y., and Wang, W. (2015), "The Emergence of GitHub as a Collaborative Platform for Education," in *Proceedings of the 18th ACM Conference on Computer Supported Cooperative Work & Social Computing, CSCW '15*, ACM, New York, NY, USA, pp. 1906–1917.